

Generalized Language Models

This PDF conversion (v1) was generated on October 10, 2025*

Online version: <https://lilianweng.github.io/posts/2019-01-31-lm/>

Date: **Estimated Reading Time:** min **Author:** Lilian Weng

⁰This file was prepared by Sakura (bili_sakura@zju.edu.cn). Please contact for any issues or copyright concerns. The original source of this file is <https://lilianweng.github.io/posts/2019-01-31-lm/>.

Contents

1	Generalized Language Models	3
2	CoVe#	5
2.1	NMT Recap#	5
2.2	Use CoVe in Downstream Tasks#	6
3	ELMo#	7
3.1	Bidirectional Language Model#	7
3.2	ELMo Representations#	8
3.3	Use ELMo in Downstream Tasks#	8
4	Cross-View Training#	8
4.1	Model Architecture#	9
4.2	Multi-Task Learning#	9
4.3	Use CVT in Downstream Tasks#	9
5	ULMFiT#	10
6	GPT#	11
6.1	Transformer Decoder as Language Model#	12
6.2	Byte Pair Encoding#	12
6.3	Supervised Fine-Tuning#	12
7	BERT#	15
7.1	Pre-training Tasks#	16
7.2	Input Embedding#	20
7.3	Use BERT in Downstream Tasks#	21
8	ALBERT#	23
8.1	Factorized Embedding Parameterization#	23
8.2	Cross-layer Parameter Sharing#	24
8.3	Sentence-Order Prediction (SOP)#	24
9	GPT-2#	24
9.1	Zero-Shot Transfer#	24
9.2	BPE on Byte Sequences#	25
9.3	Model Modifications#	25
10	RoBERTa#	25
11	T5#	26
12	GPT-3#	28
13	XLNet#	33
14	BART#	37

15 ELECTRA#	41
16 Summary#	44
17 Metric: Perplexity#	44
18 Common Tasks and Datasets#	45
19 Reference#	48
20 Citation	50
21 References	52

[Lil'Log](#)

- [|](#)
- [Posts](#)
- [Archive](#)
- [Search](#)
- [Tags](#)
- [FAQ](#)

1 Generalized Language Models

Date: January 31, 2019 | Estimated Reading Time: 36 min | Author: Lilian Weng
Table of Contents

- [CoVe](#)
 - [NMT Recap](#)
 - [Use CoVe in Downstream Tasks](#)
- [ELMo](#)
 - [Bidirectional Language Model](#)
 - [ELMo Representations](#)
 - [Use ELMo in Downstream Tasks](#)
- [Cross-View Training](#)
 - [Model Architecture](#)
 - [Multi-Task Learning](#)
 - [Use CVT in Downstream Tasks](#)
- [ULMFiT](#)
- [GPT](#)
 - [Transformer Decoder as Language Model](#)
 - [Byte Pair Encoding](#)
 - [Supervised Fine-Tuning](#)
- [BERT](#)
 - [Pre-training Tasks](#)
 - [Input Embedding](#)

- Use BERT in Downstream Tasks
- ALBERT
 - Factorized Embedding Parameterization
 - Cross-layer Parameter Sharing
 - Sentence-Order Prediction (SOP)
- GPT-2
 - Zero-Shot Transfer
 - BPE on Byte Sequences
 - Model Modifications
- RoBERTa
- T5
- GPT-3
- XLNet
- BART
- ELECTRA
- Summary
- Metric: Perplexity
- Common Tasks and Datasets
- Reference

[Updated on 2019-02-14: add [ULMFiT](#) and [GPT-2](#).]

[Updated on 2020-02-29: add [ALBERT](#).]

[Updated on 2020-10-25: add [RoBERTa](#).]

[Updated on 2020-12-13: add [T5](#).]

[Updated on 2020-12-30: add [GPT-3](#).]

[Updated on 2021-11-13: add [XLNet](#), [BART](#) and [ELECTRA](#); Also updated the [Summary](#) section.]

We have seen amazing progress in NLP in 2018. Large-scale pre-trained language models like [OpenAI GPT](#) and [BERT](#) have achieved great performance on a variety of language tasks using generic model architectures. The idea is similar to how ImageNet classification pre-training helps many vision tasks (*). Even better than vision classification pre-training, this simple and powerful approach in NLP does not require labeled data for pre-training, allowing us to experiment with increased training scale, up to our very limit.




Figure 1: I guess they are Elmo & Bert? (Image source: [here](#))

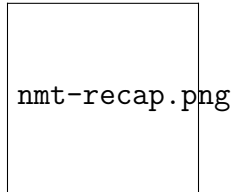


Figure 2: The NMT base model used in CoVe.

(*) *He et al. (2018) [found](#) that pre-training might not be necessary for image segmentation task.*

In my previous NLP [post on word embedding](#), the introduced embeddings are not context-specific — they are learned based on word concurrency but not sequential context. So in two sentences, “*I am eating an apple*” and “*I have an Apple phone*”, two “apple” words refer to very different things but they would still share the same word embedding vector.

Despite this, early adoption of word embeddings in problem-solving is to use them as additional features for an existing task-specific model and in a way the improvement is bounded.

In this post, we will discuss how various approaches were proposed to make embeddings dependent on context, and to make them easier and cheaper to be applied to downstream tasks in general form.

2 CoVe#

CoVe ([McCann et al. 2017](#)), short for **Contextual Word Vectors**, is a type of word embeddings learned by an encoder in an [attentional seq-to-seq](#) machine translation model. Different from traditional word embeddings introduced [here](#), CoVe word representations are functions of the entire input sentence.

2.1 NMT Recap#

Here the Neural Machine Translation (NMT) model is composed of a standard, two-layer, bidirectional LSTM encoder and an attentional two-layer unidirectional LSTM decoder. It is pre-trained on the English-German translation task. The encoder learns and optimizes the embedding vectors of English words in order to translate them to German. With the intuition that the encoder should capture high-level semantic and syntactic meanings before transforming words into another language, the encoder output is used to provide contextualized word embeddings for various downstream language tasks.




Figure 3: The CoVe embeddings are generated by an encoder trained for machine translation task. The encoder can be plugged into any downstream task-specific model. (Image source: [original paper](#))

- A sequence of n words in source language (English): $\mathbf{x} = [x_1, \dots, x_n]$.
- A sequence of m words in target language (German): $\mathbf{y} = [y_1, \dots, y_m]$.
- The GloVe vectors of source words: $\text{GloVe}(\mathbf{x})$.
- Randomly initialized embedding vectors of target words: $\mathbf{z} = [z_1, \dots, z_m]$.
- The biLSTM encoder outputs a sequence of hidden states: $\mathbf{h} = [h_1, \dots, h_n] = \text{biLSTM}(\text{GloVe}(\mathbf{x}))$ and $h_t = [\overrightarrow{h}_t; \overleftarrow{h}_t]$ where the forward LSTM computes $\overrightarrow{h}_t = \text{LSTM}(x_t, \overrightarrow{h}_{t-1})$ and the backward computation gives us $\overleftarrow{h}_t = \text{LSTM}(x_t, \overleftarrow{h}_{t-1})$.
- The attentional decoder outputs a distribution over words: $p(y_t | \mathbf{H}, y_1, \dots, y_{t-1})$ where $\mathbf{H}$ is a stack of hidden states $\{h_t\}$ along the time dimension:

$$\begin{aligned} & \text{decoder hidden state: } s_t \&= \text{LSTM}([z_{t-1}; \\ & \tilde{h}_{t-1}, s_{t-1}]) \quad \text{attention weights: } \alpha_t \&= \text{softmax}(H(W_1 \\ & s_t + b_1)) \quad \text{context-adjusted hidden state: } \tilde{h}_t \&= \tanh(W_2[H^{\top} \alpha_t; s_t] \\ & + b_2) \quad \text{decoder output: } p(y_t | \mathbf{H}, y_1, \dots, y_{t-1}) \&= \text{softmax}(W_{\text{out}} \\ & \tilde{h}_t + b_{\text{out}}) \end{aligned}$$

2.2 Use CoVe in Downstream Tasks#

The hidden states of NMT encoder are defined as **context vectors** for other language tasks:

$$\text{CoVe}(\mathbf{x}) = \text{biLSTM}(\text{GloVe}(\mathbf{x}))$$

The paper proposed to use the concatenation of GloVe and CoVe for question-answering and classification tasks. GloVe learns from the ratios of global word co-occurrences, so it has no sentence context, while CoVe is generated by processing text sequences is able to capture the contextual information.

$$\mathbf{v} = [\text{GloVe}(\mathbf{x}); \text{CoVe}(\mathbf{x})]$$

Given a downstream task, we first generate the concatenation of GloVe + CoVe vectors of input words and then feed them into the task-specific models as additional features.

Summary: The limitation of CoVe is obvious: (1) pre-training is bounded by available datasets on the supervised translation task; (2) the contribution of CoVe to the final performance is constrained by the task-specific model architecture.




Figure 4: The biLSTM base model of ELMo. (Image source: recreated based on the figure in [“Neural Networks, Types, and Functional Programming”](<http://colah.github.io/posts/2015-09-NN-Types-FP/>) by Christopher Olah.)

In the following sections, we will see that ELMo overcomes issue (1) by unsupervised pre-training and OpenAI GPT & BERT further overcome both problems by unsupervised pre-training + using generative model architecture for different downstream tasks.

3 ELMo#

ELMo, short for **E**mbeddings from **L**anguage **M**odel (Peters, et al, 2018) learns contextualized word representation by pre-training a language model in an *unsupervised* way.

3.1 Bidirectional Language Model#

The bidirectional Language Model (**biLM**) is the foundation for ELMo. While the input is a sequence of n tokens, $(x_1, \dots, x_n)$, the language model learns to predict the probability of next token given the history.

In the forward pass, the history contains words before the target token,

$$p(x_1, \dots, x_n) = \prod_{i=1}^n p(x_i \mid x_1, \dots, x_{i-1})$$

In the backward pass, the history contains words after the target token,

$$p(x_1, \dots, x_n) = \prod_{i=1}^n p(x_i \mid x_{i+1}, \dots, x_n)$$

The predictions in both directions are modeled by multi-layer LSTMs with hidden states $\overrightarrow{\mathbf{h}}_{i,\ell}$ and $\overleftarrow{\mathbf{h}}_{i,\ell}$ for input token x_i at the layer level $\ell=1, \dots, L$. The final layer’s hidden state $\mathbf{h}_{i,L} = [\overrightarrow{\mathbf{h}}_{i,L}; \overleftarrow{\mathbf{h}}_{i,L}]$ is used to output the probabilities over tokens after softmax normalization. They share the embedding layer and the softmax layer, parameterized by Θ_e and Θ_s respectively.

The model is trained to minimize the negative log likelihood (= maximize the log likelihood for true words) in both directions:

$$\begin{aligned} \mathcal{L} = - \sum_{i=1}^n \Big(& \log p(x_i \mid x_1, \dots, x_{i-1}; \Theta_e, \overrightarrow{\Theta}_{\text{LSTM}}, \Theta_s) + \\ & \log p(x_i \mid x_{i+1}, \dots, x_n; \Theta_e, \overleftarrow{\Theta}_{\text{LSTM}}, \Theta_s) \Big) \end{aligned}$$

3.2 ELMo Representations#

On top of a L -layer biLM, ELMo stacks all the hidden states across layers together by learning a task-specific linear combination. The hidden state representation for the token x_i contains $2L+1$ vectors:

$$R_i = \{ \mathbf{h}_{i,\ell} \mid \ell = 0, \dots, L \}$$

where $\mathbf{h}_{0,\ell}$ is the embedding layer output and $\mathbf{h}_{i,\ell} = [\overrightarrow{\mathbf{h}}_{i,\ell}; \overleftarrow{\mathbf{h}}_{i,\ell}]$.

The weights, s^{task} , in the linear combination are learned for each end task and normalized by softmax. The scaling factor γ^{task} is used to correct the misalignment between the distribution of biLM hidden states and the distribution of task specific representations.

$$v_i = f(R_i; \Theta^{\text{task}}) = \gamma^{\text{task}} \sum_{\ell=0}^L s^{\text{task}}_{\ell} \mathbf{h}_{i,\ell}$$

To evaluate what kind of information is captured by hidden states across different layers, ELMo is applied on semantic-intensive and syntax-intensive tasks respectively using representations in different layers of biLM:

- **Semantic task:** The *word sense disambiguation (WSD)* task emphasizes the meaning of a word given a context. The biLM top layer is better at this task than the first layer.
- **Syntax task:** The *part-of-speech (POS) tagging* task aims to infer the grammatical role of a word in one sentence. A higher accuracy can be achieved by using the biLM first layer than the top layer.

The comparison study indicates that syntactic information is better represented at lower layers while semantic information is captured by higher layers. Because different layers tend to carry different type of information, *stacking them together helps*.

3.3 Use ELMo in Downstream Tasks#

Similar to how CoVe can help different downstream tasks, ELMo embedding vectors are included in the input or lower levels of task-specific models. Moreover, for some tasks (i.e., SNLI and SQuAD, but not SRL), adding them into the output level helps too.

The improvements brought up by ELMo are largest for tasks with a small supervised dataset. With ELMo, we can also achieve similar performance with much less labeled data.

Summary: The language model pre-training is unsupervised and theoretically the pre-training can be scaled up as much as possible since the unlabeled text corpora are abundant. However, it still has the dependency on task-customized models and thus the improvement is only incremental, while searching for a good model architecture for every task remains non-trivial.

4 Cross-View Training#

In ELMo the unsupervised pre-training and task-specific learning happen for two independent models in two separate training stages. **Cross-View Training** (abbr. CVT; Clark et al., 2018) combines them into one unified semi-supervised learning procedure




Figure 5: The overview of semi-supervised language model cross-view training. (Image source: [original paper](#))

where the representation of a biLSTM encoder is improved by both supervised learning with labeled data and unsupervised learning with unlabeled data on auxiliary tasks.

4.1 Model Architecture#

The model consists of a two-layer bidirectional LSTM encoder and a primary prediction module. During training, the model is fed with labeled and unlabeled data batches alternatively.

- On *labeled examples*, all the model parameters are updated by standard supervised learning. The loss is the standard cross entropy.
- On *unlabeled examples*, the primary prediction module still can produce a “soft” target, even though we cannot know exactly how accurate they are. In a couple of auxiliary tasks, the predictor only sees and processes a restricted view of the input, such as only using encoder hidden state representation in one direction. The auxiliary task outputs are expected to match the primary prediction target for a full view of input.

In this way, the encoder is forced to distill the knowledge of the full context into partial representation. At this stage, the biLSTM encoder is backpropagated but the primary prediction module is *fixed*. The loss is to minimize the distance between auxiliary and primary predictions.

4.2 Multi-Task Learning#

When training for multiple tasks simultaneously, CVT adds several extra primary prediction models for additional tasks. They all share the same sentence representation encoder. During supervised training, once one task is randomly selected, parameters in its corresponding predictor and the representation encoder are updated. With unlabeled data samples, the encoder is optimized jointly across all the tasks by minimizing the differences between auxiliary outputs and primary prediction for every task.

The multi-task learning encourages better generality of representation and in the meantime produces a nice side-product: all-tasks-labeled examples from unlabeled data. They are precious data labels considering that cross-task labels are useful but fairly rare.

4.3 Use CVT in Downstream Tasks#

Theoretically the primary prediction module can take any form, generic or task-specific design. The examples presented in the CVT paper include both cases.




Figure 6: The sequential tagging task depends on four auxiliary prediction models, their inputs only involving hidden states in one direction: forward, backward, future and past. (Image source: [original paper](#))

In sequential tagging tasks (classification for every token) like [NER](#) or [POS](#) tagging, the predictor module contains two fully connected layers and a softmax layer on the output to produce a probability distribution over class labels. For each token $\mathbf{x}_i$, we take the corresponding hidden states in two layers, $\mathbf{h}_1^{(i)}$ and $\mathbf{h}_2^{(i)}$:

$$\begin{aligned} p_{\theta}(y_i \mid \mathbf{x}_i) &= \text{NN}(\mathbf{h}_1^{(i)}) \\ &= \text{NN}([\mathbf{h}_1^{(i)}; \mathbf{h}_2^{(i)}]) \\ &= \text{softmax}(\mathbf{W} \cdot \text{ReLU}(\mathbf{W}' \cdot [\mathbf{h}_1^{(i)}; \mathbf{h}_2^{(i)}]) + \mathbf{b}) \end{aligned}$$

The auxiliary tasks are only fed with forward or backward LSTM state in the first layer. Because they only observe partial context, either on the left or right, they have to learn like a language model, trying to predict the next token given the context. The **fwd** and **bwd** auxiliary tasks only take one direction. The **future** and **past** tasks take one step further in forward and backward direction, respectively.

$$\begin{aligned} p_{\theta}^{\text{fwd}}(y_i \mid \mathbf{x}_i) &= \text{NN}^{\text{fwd}}(\overrightarrow{\mathbf{h}_1}) \\ p_{\theta}^{\text{bwd}}(y_i \mid \mathbf{x}_i) &= \text{NN}^{\text{bwd}}(\overleftarrow{\mathbf{h}_1}) \\ p_{\theta}^{\text{future}}(y_i \mid \mathbf{x}_i) &= \text{NN}^{\text{future}}(\overrightarrow{\mathbf{h}_1}) \\ p_{\theta}^{\text{past}}(y_i \mid \mathbf{x}_i) &= \text{NN}^{\text{past}}(\overleftarrow{\mathbf{h}_1}) \end{aligned}$$

Note that if the primary prediction module has dropout, the dropout layer works as usual when training with labeled data, but it is not applied when generating “soft” target for auxiliary tasks during training with unlabeled data.

In the machine translation task, the primary prediction module is replaced with a standard unidirectional LSTM decoder with attention. There are two auxiliary tasks: (1) apply dropout on the attention weight vector by randomly zeroing out some values; (2) predict the future word in the target sequence. The primary prediction for auxiliary tasks to match is the best predicted target sequence produced by running the fixed primary decoder on the input sequence with [beam search](#).

5 ULMFiT#

The idea of using generative pretrained LM + task-specific fine-tuning was first explored in ULMFiT ([Howard & Ruder, 2018](#)), directly motivated by the success of using ImageNet pre-training for computer vision tasks. The base model is [AWD-LSTM](#).

ULMFiT follows three steps to achieve good transfer learning results on downstream language classification tasks:

1. *General LM pre-training*: on Wikipedia text.




Figure 7: Three training stages of ULMFiT. (Image source: [original paper](#))

2. *Target task LM fine-tuning*: ULMFiT proposed two training techniques for stabilizing the fine-tuning process. See below.
 - **Discriminative fine-tuning** is motivated by the fact that different layers of LM capture different types of information (see [discussion](#) above). ULMFiT proposed to tune each layer with different learning rates, $\{\eta^1, \dots, \eta^{\ell}, \dots, \eta^L\}$, where η is the base learning rate for the first layer, η^{ℓ} is for the ℓ -th layer and there are L layers in total.
 - **Slanted triangular learning rates (STLR)** refer to a special learning rate scheduling that first linearly increases the learning rate and then linearly decays it. The increase stage is short so that the model can converge to a parameter space suitable for the task fast, while the decay period is long allowing for better fine-tuning.
3. *Target task classifier fine-tuning*: The pretrained LM is augmented with two standard feed-forward layers and a softmax normalization at the end to predict a target label distribution.
 - **Concat pooling** extracts max-polling and mean-polling over the history of hidden states and concatenates them with the final hidden state.
 - **Gradual unfreezing** helps to avoid catastrophic forgetting by gradually unfreezing the model layers starting from the last one. First the last layer is unfrozen and fine-tuned for one epoch. Then the next lower layer is unfrozen. This process is repeated until all the layers are tuned.

6 GPT#

Following the similar idea of ELMo, OpenAI **GPT**, short for **Generative Pre-training Transformer** ([Radford et al., 2018](#)), expands the unsupervised language model to a much larger scale by training on a giant collection of free text corpora. Despite of the similarity, GPT has two major differences from ELMo.

1. The model architectures are different: ELMo uses a shallow concatenation of independently trained left-to-right and right-to-left multi-layer LSTMs, while GPT is a multi-layer transformer decoder.
2. The use of contextualized embeddings in downstream tasks are different: ELMo feeds embeddings into models customized for specific tasks as additional features, while GPT fine-tunes the same base model for all end tasks.




Figure 8: The transformer decoder model architecture in OpenAI GPT.

6.1 Transformer Decoder as Language Model#

Compared to the [original transformer](#) architecture, the [transformer decoder](#) model discards the encoder part, so there is only one single input sentence rather than two separate source and target sequences.

This model applies multiple transformer blocks over the embeddings of input sequences. Each block contains a masked *multi-headed self-attention* layer and a *pointwise feed-forward* layer. The final output produces a distribution over target tokens after softmax normalization.

The loss is the negative log-likelihood, same as [ELMo](#), but without backward computation. Let's say, the context window of the size k is located before the target word and the loss would look like:

$$\mathcal{L}_{\text{LM}} = -\sum_i \log p(x_i \mid x_{i-k}, \dots, x_{i-1})$$

6.2 Byte Pair Encoding#

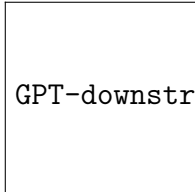
Byte Pair Encoding (BPE) is used to encode the input sequences. BPE was originally proposed as a data compression algorithm in 1990s and then was adopted to solve the open-vocabulary issue in machine translation, as we can easily run into rare and unknown words when translating into a new language. Motivated by the intuition that rare and unknown words can often be decomposed into multiple subwords, BPE finds the best word segmentation by iteratively and greedily merging frequent pairs of characters.

6.3 Supervised Fine-Tuning#

The most substantial upgrade that OpenAI GPT proposed is to get rid of the task-specific model and use the pre-trained language model directly!

Let's take classification as an example. Say, in the labeled dataset, each input has n tokens, $\mathbf{x} = (x_1, \dots, x_n)$, and one label y . GPT first processes the input sequence $\mathbf{x}$ through the pre-trained transformer decoder and the last layer output for the last token x_n is $\mathbf{h}_L^{(n)}$. Then with only one new trainable weight matrix $\mathbf{W}_y$, it can predict a distribution over class labels.

GPT-classification.png					
e-3e-3	e-3e-3	e-3e-3	e-3e-3	e-3e-3	e-3e-3
	$P(y \mid x_1, \dots, x_n)$	The loss is to minimize $-\sum_{i=1}^n \log P(y \mid x_1, \dots, x_n)$	•	$\begin{aligned} (1) \quad & \begin{aligned} & \text{it helps achieve a better structure for training and} \\ & \text{the negative log-likelihood for true labels.} \end{aligned} \\ & \text{In addition, adding the LM loss as an auxiliary loss is found to be beneficial, because:} \end{aligned}$	$\begin{aligned} & \text{With} \\ & \text{similarity} \\ & \text{lar} \\ & \text{signs,} \\ & \text{no} \\ & \text{cus-} \\ & \text{tomized} \\ & \text{model} \\ & \text{struc-} \\ & \text{ture} \\ & \text{is} \\ & \text{needed} \\ & \text{for} \\ & \text{other} \\ & \text{end} \\ & \text{tasks} \\ & \text{(see} \\ & \text{Fig. 7).} \\ & \text{the} \\ & \text{task} \\ & \text{in-} \\ & \text{put} \\ & \text{con-} \\ & \text{tains} \\ & \text{mul-} \\ & \text{ti-} \\ & \text{ple} \\ & \text{sen-} \\ & \text{tences,} \\ & \text{a} \\ & \text{spe-} \\ & \text{cial} \\ & \text{de-} \\ & \text{lim-} \\ & \text{iter} \\ & \text{to-} \\ & \text{ken} \\ & \text{(\$)} \\ & \text{is} \\ & \text{added} \\ & \text{be-} \\ & \text{tween} \\ & \text{each} \\ & \text{pair} \\ & \text{of} \end{aligned}$



GPT-downstream-tasks.png

e-3e-3For [t]5.0e-3 Training objects in slightly modified GPT transformer

the
sen-
tence
sim-
i-
lar-
ity
task,
be-
cause
the
or-
der-
ing
does
not
mat-
ter,
both
or-
der-
ings
are
in-
cluded.
For
the
mul-
ti-
ple
choice
task,
the
con-
text
is
paired
with
ev-
ery
an-
swer
can-
di-
date.


models for downstream tasks. (Image source: [original paper](#))

Summary

It is a summary of the paper. The paper is a survey of the state-of-the-art in natural language processing. It discusses the limitations of current models and the need for a more general model. The paper introduces a new model, which is a generalization of the current state-of-the-art. The model is trained on a large dataset of text and is able to perform a wide range of tasks, including machine translation, text classification, and question answering. The model is trained using a self-supervised learning objective, which allows it to learn from a large amount of unlabeled data. The paper shows that the new model outperforms the current state-of-the-art on a wide range of tasks. The paper also discusses the limitations of the new model and the need for further research.

7

BERT


 transformer-encoder-2.png

e-3e-3 BERT, 3 Comp e-3e-3 e-3e-3 The [5.0e-3] Recap of Transformer

short to “bidirectional model”

for GPT, nature of our model is the single most im-

Bidi- the largest difference and im-

rec- dif- fer- model of BERT is a multi-

tional fer- model of BERT is a multi-

En- ence is the is a multi-

coder and the is a multi-

Rep- im- prove- ment of im- rec-

re- prove- ment of im- rec-

sen- ment of im- rec-

ta- of im- rec-

tions BERT is new Trans-

from is new Trans-

Trans- to make train- ing bi-

form- make train- ing bi-

ers train- ing bi-

(De- vin, et al., 2019)

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

is model learns to pre- dict both con- text on the left and right. The pa- per ac- cord- ing to the ab- la- tion study claimed that:

Encoder model architecture. (Image source: [Transformer paper](#))

Task 3

1: Mask language model (MLM)

en- cour- age the bi- direction- pre- dic- tion and sentence- level un- der- stand- ing, BERT is trained with two tasks in- stead of the ba- sic lan- guage task (that is, to pre- dict the next to- ken given con- text).

From [Wikipedia](#): “A cloze test (also cloze deletion test) is an exer- cise, test, or as- sess- ment con- sist- ing of a por- tion of lan- guage with cer- tain items, words, or signs re- moved (cloze text), where the par- tic- i- pant is asked to re- place

language-model-comparison.png

e-3e-3	It e-3e-3	e-3e-3	e-3e-3	Motivated	e-3e-3	The	5.0e-3	Comparison
is	un-sur-pris-ing to be-lieve that a rep-re-sen-ta-tion that learns the con-text around a word rather than just af-ter the word is able to bet-ter cap-ture its mean-ing, both syn-tac-ti-cally and se-man-	1. Task Randomly mask 15% of tokens in each sequence. Because if we only re-place masked tokens with a spe-cial place-holder [MASK], the spe-cial to-ken would never be en-coun-tered dur-ing fine-tuning. Hence, BERT employed sev-	2. Next sentence pre-dic-tion by the fact that many down-stream tasks in-volve the un-der-stand-ing of re-la-tion-ships be-tween sen-tences (i.e., QA, NLI), BERT added an-other aux-il-iary task on train-ing a <i>bi-nary clas-si-fier</i> for telling whether one	3. Motivated by the fact that many down-stream tasks in-volve the un-der-stand-ing of re-la-tion-ships be-tween sen-tences (i.e., QA, NLI), BERT added an-other aux-il-iary task on train-ing a <i>bi-nary clas-si-fier</i> for telling whether one	1. Training data for both (A, B) auxiliary tasks Sample ing sen-tence pairs (A, B) so that: • above can be 50% of the gen-time, er B ated fol-lows any A; mono-fin-gual 50% of the time, Hence the scale of train-ing is low un-bounded. 2. The train-ing pro-cesses both the sen-tences and the out-put a LM bi-like-nary la-hood	1. Training data for both (A, B) auxiliary tasks Sample ing sen-tence pairs (A, B) so that: • above can be 50% of the gen-time, er B ated fol-lows any A; mono-fin-gual 50% of the time, Hence the scale of train-ing is low un-bounded. 2. The train-ing pro-cesses both the sen-tences and the out-put a LM bi-like-nary la-hood	5.0e-3	Comparison

of BERT, OpenAI GPT and ELMo model architectures. (Image source: [original paper](#))



source: [original paper](#))

e-3e-3Note e-3e-3

7.3 Use

e-3e-3BERT e-3e-3Fore-3e-3Fore-3e-3Over[110e-3

that
the
first
to-
ken
is
al-
ways
forced
to
be
[CLS]
—

a
place-
holder
that
will
be
used
later
for
pre-
dic-
tion
in
down-
stream
tasks.

the
BERT
in
Down-
stream
Tasks#
new
pa-
ram-
e-
ters
added,
just
like
Ope-
nAI
GPT.

clas-
si-
fi-
ca-
tion
tasks,
we
get
the
pre-
dic-
tion
by
tak-
ing
the
fi-
nal
hid-
den
state
of
the
spe-
cial
first
to-
ken
[CLS],
 $\text{softmax}(\text{W}_L \cdot \text{h}^{\text{CLS}})$
and
mul-
ti-
ply-
ing
it
with
a
small
weight
ma-
trix,
 $\text{softmax}(\text{W}_L \cdot \text{h}^{\text{CLS}})$
 $\text{softmax}(\text{W}_L \cdot \text{h}^{\text{CLS}})$

QA
tasks
like
SQuAD, part
we
need
to
pre-
dict
the
text
span
in
the
given
para-
graph
for
an
given
ques-
tion.
BERT
pre-
dicts
two
prob-
a-
bil-
ity
dis-
tri-
bu-
tions
of
ev-
ery
to-
ken,
be-
ing
the
pa-
ram-
eters
for
the
im-
ple-
men-
ta-
tion
de-
tails
for

the
add-
on
for
end
task
fine-
tuning
is
very
min-
imal
—
one
or
two
weight
ma-
tri-
ces
to
con-
vert
the
Trans-
form
den
states
to
an
in-
ter-
pretable
for-
mat.
Check
the
pa-
ram-
eters
for
the
im-
ple-
men-
ta-
tion
de-
tails
for



BERT-downstream-tasks.png

Training objects in slightly modified BERT models for downstream tasks.

(Image source: [original paper](#)) e-3e-3A e-3e-3| e-3e-3 e-3e-3ALBERTe-3e-3Ine-3e-3Conceptually,

summary table compares differences between fine-tuning of OpenAI GPT and BERT.	OpenAI GPT BERT Special character [SEP] and [CLS] are only introduced at fine-tuning stage. [SEP] and [CLS] and sentence A/B embeddings are learned at the pre-training stage. Training process 1M steps, batch	8	ALBERT (Lai et al. 2019), short for A Lite BERT , is a light-weighted version of BERT model. An ALBERT model can be trained 1.7x faster with 18x fewer parameters, compared to a BERT model of similar con-fig-uration. ALBERT in-	Factorized BERT, i.e. Embedding Piecewise Parameterization of embedding is ex-pected to learn <i>context-independent</i> representations and the hidden states are <i>context-dependent</i> , it makes sense to separate the size of the hidden layers from the size of the vocabulary-
--	---	---	---	---

8.2	Cross-layer parameterization	8.3	Sentence Order Prediction (SOP) (NSP)	9	GPT-2	9.1	Zero-Shot Transfer for GPT-2
	layer parameters can be shared in many ways: (a) only share feed-forward part; (b) only share attention parameters; or (c) share all the parameters. This technique reduces the number of parameters by a ton	the Order Prediction (SOP) task of BERT turned out to be too easy. AL-BERT instead adopted a sentence-order prediction (SOP) self-supervised loss,	Positive NSP task, the model can make a decision if it is able to detect top-ones when above, but the segments from different texts. In comparison, SOP is harder as it requires the model to fully understand		GPT-2 is a language model is a direct successor to GPT-2 has 1.5B parameters, 10x more than the original GPT, and it achieves SOTA results on 7 out of 8 tested language modeling datasets in a	Zero-Shot Transfer for GPT-2 is solely a language model. All the downstream language tasks are framed as predicting conditional probabilities and there is no task-specific fine-tuning.	

e-3e-3 e-3e-3 Same e-3e-3 BPE e-3e-3 Using e-3e-3 e-3e-3 Compared e-3e-3 e-3e-3 RoBERTa e-3e-3 RoBERTa

9.2

BPE

as merges the on the fre- byte orig- quently se- Byte co- quence Se- nal occurred rep- quences# GPT- pairs sen- 2 in ta- uses a tion, BPE greedy GPT- but man- 2 is UTF- To able 8 pre- to byte vent as- se- it sign quences from a Each gen- prob- byte er- a- can at- bil- rep- ing ity re- re- mul- to 256 ple Uni- dif- ver- code fer- sions string, ent of re- val- com- gard- ues mon less in words of 8 (i.e. any bits, dog., pre- while dog! processing UTF- and steps. 8 dog? can for use the word up dog), GPT- 4 GPT- 2 bytes pre- for one vents char- BPE ac- from merg- sup- ing char- port- ac- ing up ters across to sat-

9.3

Model

No MBD- i- other than hav- fi- cag- tions# more trans- former lay- ers and pa- ram- e- ters, GPT- 2 in- cor- po- rates only a few ar- chi- tec- ture mod- i- fi- ca- tions:

Layer

nor- mal- iza- tion was moved to the in- put of each sub- block, sim- i- lar to a resid- ual unit of type “build- ing block” (dif- fer- ently from the orig- i- nal type “bot- tle- neck”, it has batch nor- mal- iza- tion ap- plied be- fore weight

RoBERTa

for Robustly optimized BERT approach; Liu, et al. 2019) refers to a new receipt for train- ing BERT to achieve bet- ter re- sults, as they found that the orig- i- nal BERT model is sig- nif- i- cantly un- der- trained. The re- ceipt con- tains the fol- low-

2.

Train- added for a longer new dataset with big- ger batch size. and Remove the next sen- tence pre- dic- tion (NSP) task. more data helps. Use longer se- quences in for- train- ing on down- stream tasks. The pawas pretrained with the us BPE ing on in-byte di-se- quences, ual same sen- tences as GPT- in-2. They also found that per-


e-3e-3

11

e-3e-3

T5

[t]5.0e-3



A diagram of T5 task evaluation. The

language model

T5

is short for “Text-to-Text Transfer Transformer” (Rafel et al., 2020). The encoder-decoder implementation follows the original Trans-former architecture: token embedding → encoder

text-to-text framework casts every task into a generic form: feeding input text to predict

some target text. (Image source: [Raffel et al., 2020](#))

e-3e-3The e-3e-3As e-3e-3

12

GPT

3#

model the
is au-
trained thors
on men-
Web tioned
cor- in
pus the
ex- pa-
tracted per
from "... our
Apr goal
2019 is
with not
var- to
i- pro-
ous pose
fil- new
ters meth-
ap- ods
plied. but
The in-
model stead
is to
fine- pro-
tuned vide
for a
each com-
down- pre-
stream hen-
task sive
sep- per-
a- spec-
rately tive
via on
"adapter where
lay- the
ers" field
(add stands",
an the
ex- T5
tra long
layer pa-
for per
train- de-
ing) scribed
or a
"grad- lot
ual of
un- train-
freez- ing
ing" setup
(see and

e-3e-3GPT-

3

(Brown

et

al.,

2020)

has

the

same

ar-

chi-

tec-

ture

as

GPT-

2

but

con-

tains

175B

pa-

ram-

e-

ters,

10x

larger

than

GPT-

2

(1.5B).

In

ad-

di-

tion,

GPT-

3

uses

al-

ter-

nat-

ing

dense

and

lo-

cally


banded

sparse

at-

ten-

[t]5.0e-3



GPT3-train-data.png

Training datasets for

GPT-3. Note that the occurrence of each dataset during training is not proportional to

A rectangular box containing the text "GPT3-eval.png".

the dataset size. (Table source: [Brown et al., 2020](#))


e-3e-3For [t]5.0e-3

all
the
down-
stream
eval-
u-
a-
tion,
GPT-
3
is
tested
in
the
few-
shot
set-
ting
with-
out
any
gradient-
based
fine-
tuning.
Here
the
few-
shot
ex-
am-
ples
are
pro-
vided
as
part
of
the
prompt.
GPT-
3
achieves
strong
per-
for-
mance
on

The evaluation performance increases with the model size and the number of examples.

(Image source: [Brown et al., 2020](#)) e-3e-3 e-3e-3The-3XLNet-3e-3where-3Not-3However,

13 **XLNet** $\backslash\begin{\mathbb{Z}}-\text{Two}$
 tore- **et** $\backslash\mathcal{L}-\text{text}\{\text{XLNet}\}$
 gres- **al.** $\&=$ a naive fer-
 sive **2019)** - set rep-ent
 (*AR*) gen- $\backslash\mathbf{of}\{E\}-\{\backslash\mathbf{of}\{z\}$
 model er- $\sim$ all sen-quire-
 such al- $\backslash\mathcal{H}\{\mathbb{Z}\}-\text{Ta-}$ ments
 as izes $\backslash\text{Big[}$ si- tion on
 GPT the $\backslash\sum_{\{t\}}^{\text{To}}$ $\$g_{\backslash\theta}$
 and AE $\backslash\log$ per- the $(\backslash\mathbf{of}\{x\}-\{\backslash$
au- method $p_{\backslash\theta}$ mu- hid- $z_t)\$$
toen- to $(X_{\{z,t\}}$ a- den lead
coder in- $=$ tion state to
 (*AE*) cor- x of of a
 model po- $\backslash\text{mid}$ length the two-
 such rate $\backslash\mathbf{of}\{E\}-\{\backslash\mathbf{of}\{z\}\text{year}\{\text{t}\}\}\backslash\text{I}$
 as the $\backslash\backslash$ $\$z_t$ text, self-
 BERT ben- $\&=$ and $\$h_{\backslash\theta}$ attention
 are e- - $\$ \backslash\mathbf{of}\{z\}\text{year}\{\text{t}\}\backslash\{\backslash\mathbf{of}\{x\}$
 two fits $\backslash\mathbf{of}\{E\}-\{\backslash\mathbf{of}\{z\}$
 most of $\sim$ note red, to
 com- *AR.* $\backslash\mathcal{H}\{\mathbb{Z}\}-\text{To}$ does ac-
 mon XL- $\backslash\text{Big[}$ $\$t$ - not com-
 ways Net $\backslash\log$ th de- mo-
 for pro- $\backslash\text{frac}\{\text{el-}$ pend date:
 lan- posed $\backslash\exp(e(x)^{\backslash\text{topon}}$
 guage the $\backslash\text{color}\{\text{red}\}\{h_{\backslash\theta}\}$ which
 mod- **per-** $(\backslash\mathbf{of}\{x\}-\{\backslash\mathbf{of}\{z\}-\{<\{t\}\}\})$
 el- **mu-** $\}\{\text{the si-}$
 ing. **ta-** $\backslash\sum_{\{t\}}$ tion
 How- **tion** $\backslash\exp(e(x)^{\backslash\text{top}}$ the
 ever, **lan-** $\backslash\text{color}\{\text{red}\}\{h_{\backslash\theta}\}$ which
 each **guage** $(\backslash\mathbf{of}\{x\}-\{\backslash\mathbf{of}\{z\}-\{<\{t\}\}\})$
 has **mod-** $\}$ e- to
 their **el-** $\backslash\text{Big[}$ ments pre-
 own **ing** $\backslash\backslash$ of dict,
 dis- ob- $\&=$ a as
 ad- jec- - per- the
 van- tive. $\backslash\mathbf{of}\{E\}-\{\backslash\mathbf{of}\{z\}$
 tages: For $\sim$ ta- mu-
 AR a $\backslash\mathcal{H}\{\mathbb{Z}\}-\text{Ta-}$
 does text $\backslash\text{Big[}$ $\$ \backslash\mathbf{of}\{z\}$
 not se- $\backslash\log$ $\backslash\text{in}$ breaks
 learn quence, $\backslash\text{frac}\{\backslash\mathcal{H}\{\mathbb{Z}\}-\text{T}\$$.
 the it $\backslash\exp(e(x)^{\backslash\text{topde-}}$
 bidi- sam- $\backslash\text{color}\{\text{blue}\}\{g_{\backslash\theta}$ $\text{f}\{\theta\}$
 rec- ples $(\backslash\mathbf{of}\{x\}-\{\backslash\mathbf{of}\{z\}-\{<\{t\}\}\},$
 tional a $z_t))$ der-
 con- fac- $\}\{\text{ing.}$
 text, tor- $\backslash\sum_{\{x'\}}$ There-
 which iza- $\backslash\exp(e(x')^{\backslash\text{top}}$ fore,
 is tion $\backslash\text{color}\{\text{blue}\}\{g_{\backslash\theta}$ $\text{f}\{\theta\}$


 XLNet-two-stream-attention.png

e-3e-3 [t]5.0e-3 1. The illustration of two-stream self-attention mechanism


When
 pre-
 dict-
 ing
 x_{z_t} ,
 it
 should
 only
 en-
 code
 the
 po-
 si-
 tion
 z_t
 but
 not
 the
 con-
 tent
 x_{z_t} ;
 oth-
 er-
 wise
 it
 is
 triv-
 ial.
 This
 is
 wrapped
 into
 the
 “query
 rep-
 re-
 sen-
 ta-
 tion”
 g_{z_t}
 =
 $g_{\theta}(\mathbf{x}_{\mathbf{z}_{<t}})$,
 z_t
 does
 not
 en-

in XLNet. (Image source: [Yang et al. 2019](#))

Conceptually, Given the two streams of representations are updated as follows,

$$\begin{aligned} \mathbf{h}_{z,t}^{(m)} &= \mathbf{h}_{z,t}^{(m-1)} + \mathbf{g}_{z,t}^{(m)} \mathbf{K} \mathbf{V}^T \mathbf{h}_{z,t}^{(m-1)} \\ \mathbf{g}_{z,t}^{(m)} &= \text{softmax}(\mathbf{h}_{z,t}^{(m-1)} \mathbf{W}_g \mathbf{h}_{z,t}^{(m-1)}) \end{aligned}$$

where $\mathbf{h}_{z,t}^{(m)}$ is the hidden state at time t for the m -th stream, $\mathbf{g}_{z,t}^{(m)}$ is the gating vector, $\mathbf{K}$ and $\mathbf{V}$ are the key and value matrices, and $\mathbf{W}_g$ is the gating weight matrix. The hidden state $\mathbf{h}_{z,t}^{(m)}$ is updated by adding the product of the gating vector and the key-value product to the previous hidden state.



Yang et al. 2019)

e-3e-3

14

BART

e-3e-3

BART

#

(Lewis

et

al.,

2019)

is

a

de-

nois-

ing

au-

toen-

coder

to

re-

cover

the

orig-

i-

nal

text

from

a

ran-

domly

cor-

rupted

ver-

sion.

It

com-

bines

Bidirectional

and

AutoRegressive

Transformer:

pre-

cisely,

jointly

train-

ing

BERT-

like

bidi-

rec-

tional

en-

coder

and


GPT-

like

au-

tore-

[t]5.0e-3




BART.png

A schematic comparison of BART with BERT and GPT. (Image source:

[Lewis et al., 2019](#))

e-3e-3The[xt]5.0e-3



BART-perf.png

Comparison of different language

ex-
per-
i-
mented
with
a
va-
ri-
ety
of
nois-
ing
trans-
for-
ma-
tions,
in-
clud-
ing
to-
ken
mask-
ing,
to-
ken
dele-
tion,
text
in-
fill-
ing
(i.e.
A
ran-
domly
sam-
pled
text
span,
which
may
con-
tain
mul-
ti-
ple
to-
kens,

modeling pre-training objectives. (Image source: [Lewis et al., 2019](#))

e-3e-3 Learning

from
their
ex-
per-
i-
ments:

- The performance of pre-training methods varies significantly across downstream tasks.

- Token masking is crucial, as the performance is poor when only sentence permutation or document permutation is applied.


 ELECTRA-overview.png

e-3e-3 e-3e-3 Most 5.0e-3

Illustration of ELECTRA model architecture.


15 ELECTRA#

rent
 pre-
 training
 large
 lan-
 guage
 mod-
 els
 de-
 mand
 a
 lot
 of
 com-
 pu-
 ta-
 tion
 re-
 sources,
 rais-
 ing
 con-
 cerns
 about
 their
 cost
 and
 ac-
 ces-
 si-
 bil-
 ity.
ELEC-
TRA
 (“Ef-
 fi-
 ciently
 Learn-
 ing
 an
 En-
 coder
 that
 Clas-
 si-
 fies
 To-

(Image source: [Clark et al. 2020](#)) e-3e-3ELBC-TFA-3e-3The-3\$e-3The-3After-3

pro-
poses
a
new
pre-
train-
ing
task,
called
“Re-
placed
To-
ken
De-
tec-
tion”
(RTD).
Let’s
ran-
domly
sam-
ple
\$k\$
po-
si-
tions
to
be
masked.
Each
se-
lected
to-
ken
in
the
orig-
inal
text
is
re-
placed
by
a
plau-
si-
ble
al-
ter-
na-
tive

$\begin{aligned} & \backslash \text{begin}\{\text{aligned}\} \backslash \text{begin}\{\text{aligned}\} \text{pre-} \\ & \text{pro-} \quad \backslash \text{boldsymbol}\{\text{m}\} \text{mathcal}\{L\} \backslash \text{text}\{\text{MLM}\} (\backslash \text{mathbf}\{x\} \\ & \text{a} \quad \&= \quad \text{the} \quad \backslash \text{theta}_{\text{G}} \text{more ing} \\ & \text{new} \quad [\text{m}_1, \text{gen-} \quad \&= \quad \text{ben-} \quad \text{the} \\ & \text{pre-} \quad \backslash \text{dots}, \text{er-} \quad \backslash \text{mathbb}\{E\} \backslash \text{Big}(\backslash \text{sum}_{\text{i}} \\ & \text{train-} \quad \text{m}_k] \text{a-} \quad \backslash \text{in} \quad \text{fi-} \quad \text{er-} \\ & \text{ing} \quad \backslash \text{text}\{\text{tor} \quad \backslash \text{boldsymbol}\{\text{m}\} \\ & \text{task,} \quad \text{where} \quad \text{is} \quad - \quad \text{to} \quad \text{tor} \\ & \text{called} \quad \} \quad \text{the} \quad \backslash \text{log} \quad \text{only} \quad \text{is} \\ & \text{“Re-} \quad \text{m}_i \quad \text{neg-} \quad \text{p}_G \quad \text{share} \quad \text{dis-} \\ & \text{placed} \quad \backslash \text{sim} \quad \text{a-} \quad (\text{x}_i \quad \text{the} \quad \text{carded} \\ & \text{To-} \quad \backslash \text{text}\{\text{until}\} \backslash \{1, \backslash \text{mid} \quad \text{em-} \quad \text{and} \\ & \text{ken} \quad \text{n}\} \backslash \text{text}\{\text{log-} \quad \backslash \text{boldsymbol}\{\text{m}\} \backslash \text{only} \backslash \text{text}\{\text{masked}\} \\ & \text{De-} \quad \text{for} \quad \text{likelihood}(\text{Big}) \quad \text{dings} \quad \text{the} \\ & \text{tec-} \quad \} \quad \text{just} \quad \backslash \backslash \quad \text{be-} \quad \text{ELEC-} \\ & \text{tion”} \quad \text{i}=1, \quad \text{as} \quad \backslash \text{mathcal}\{\text{Dep}\} \backslash \text{Text}\{\text{Disc}\} (\backslash \text{mathbf}\{x\}, \\ & \text{(RTD).} \quad \backslash \text{dots}, \quad \text{in} \quad \backslash \text{theta}_{\text{D}} \text{on-} \quad \text{dis-} \\ & \text{Let’s} \quad \text{k} \quad \text{other} \quad \&= \quad \text{er-} \quad \text{crim-} \\ & \text{ran-} \quad \backslash \backslash \quad \text{lan-} \quad \backslash \text{mathbb}\{E\} \backslash \text{Big}(\backslash \text{sum}_{\text{i}} \\ & \text{domly} \quad \backslash \text{boldsymbol}\{\text{m}\} \backslash \text{only} \backslash \text{text}\{\text{masked}\} \backslash \text{h-} \\ & \text{sam-} \quad \&= \quad \text{mod-} \quad \backslash \text{mathbb}\{E\} \backslash \text{Big}(\backslash \text{sum}_{\text{i}} \backslash \text{text}\{\text{corrupt}\} \backslash \text{t} \\ & \text{ple} \quad \backslash \text{text}\{\text{REPLACE}\} (\backslash \text{boldsymbol}\{\text{m}\} \backslash \text{only} \backslash \text{text}\{\text{masked}\} \backslash \text{h-} \\ & \k \quad \backslash \text{boldsymbol}\{\text{m}\} \backslash \text{only} \backslash \text{text}\{\text{masked}\} \backslash \text{h-} \quad \text{crim-} \quad \text{fine-} \\ & \text{po-} \quad \backslash \text{texttt}\{[MASK]\} \backslash \text{log} \quad \text{i-} \quad \text{tuned} \\ & \text{si-} \quad \backslash \backslash \quad \text{for} \quad \text{D}(\backslash \text{boldsymbol}\{\text{m}\} \backslash \text{only} \backslash \text{text}\{\text{masked}\} \backslash \text{h-} \\ & \text{tions} \quad \backslash \text{boldsymbol}\{\text{m}\} \backslash \text{only} \backslash \text{text}\{\text{masked}\} \backslash \text{h-} \quad \text{ther} \\ & \text{to} \quad \&= \quad \text{dis-} \quad - \quad \text{while} \quad \text{for} \\ & \text{be} \quad \backslash \text{text}\{\text{REPLACE}\} (\backslash \text{boldsymbol}\{\text{m}\} \backslash \text{only} \backslash \text{text}\{\text{masked}\} \backslash \text{h-} \\ & \text{masked.} \quad \backslash \text{boldsymbol}\{\text{m}\} \backslash \text{only} \backslash \text{text}\{\text{masked}\} \backslash \text{h-} \quad \text{ing} \quad \text{stream} \\ & \text{Each} \quad \backslash \text{tilde}\{\backslash \text{boldsymbol}\{\text{m}\} \backslash \text{only} \backslash \text{text}\{\text{masked}\} \backslash \text{h-}\} \quad \text{tasks.} \\ & \text{se-} \quad \backslash \text{text}\{\text{tor} \quad \backslash \text{log} \quad \text{small} \quad \text{The} \\ & \text{lected} \quad \text{where} \quad \text{is} \quad (1 \quad \text{gen-} \quad \text{fol-} \\ & \text{to-} \quad \} \quad \text{the} \quad - \quad \text{er-} \quad \text{low-} \\ & \text{ken} \quad \backslash \text{tilde}\{\text{x}\} \backslash \text{cross-} \quad \backslash \text{log} \quad \text{a-} \quad \text{ing} \\ & \text{in} \quad \backslash \text{sim} \quad \text{entropyD}(\backslash \text{boldsymbol}\{\text{m}\} \backslash \text{only} \backslash \text{text}\{\text{masked}\} \backslash \text{h-} \\ & \text{the} \quad \text{p}_G(\text{x}_i \text{Note} \quad \text{t})) \quad (1/4 \quad \text{ble} \\ & \text{orig-} \quad \backslash \text{mid} \quad \text{that} \quad \backslash \text{Big}) \quad \text{to} \quad \text{shows} \\ & \text{i-} \quad \backslash \text{boldsymbol}\{\text{m}\} \backslash \text{only} \backslash \text{text}\{\text{masked}\} \backslash \text{h-} \quad \text{ELEC-} \\ & \text{nal} \quad \backslash \text{text}\{\text{gen-} \quad \$\$ \quad \text{the} \quad \text{TRA’s} \\ & \text{text} \quad \text{for} \quad \text{er-} \quad \text{dis-} \quad \text{per-} \\ & \text{is} \quad \} \quad \text{a-} \quad \text{crim-} \quad \text{for-} \\ & \text{re-} \quad \text{i} \quad \text{tor} \quad \text{i-} \quad \text{mance} \\ & \text{placed} \quad \backslash \text{in} \quad \text{is} \quad \text{na-} \quad \text{on} \\ & \text{by} \quad \backslash \text{boldsymbol}\{\text{m}\} \backslash \text{only} \backslash \text{text}\{\text{masked}\} \backslash \text{h-} \quad \text{tor} \quad \text{the} \\ & \text{a} \quad \backslash \backslash \quad \text{ad-} \quad \text{size),} \quad \text{GLUE} \\ & \text{plau-} \quad \backslash \text{end}\{\text{aligned}\} \quad \text{rather} \quad \text{dev} \\ & \text{si-} \quad \$\$ \quad \text{sar-} \quad \text{than} \quad \text{set.} \\ & \text{ble} \quad \text{i-} \quad \text{shar-} \\ & \text{al-} \quad \text{ally} \quad \text{ing} \\ & \text{ter-} \quad \text{trained} \quad \text{all} \\ & \text{na-} \quad \text{to} \quad \text{the} \\ & \text{tive} \quad \text{fool} \quad \text{weights} \end{aligned}$



ELECTRA-perf.png

Comparison of ELECTRA with other language models on the GLUE dev

e-3e-3	Given	e-3e-3	The	e-3e-3	A	e-3e-3	e-3e-3	e-3e-3	e-3e-3	e-3e-3	Commonsense	e-3e-3	Natural										
a	H(s)	per-	\begin{aligned}	18	Commonsense	Tasks	and	Datasets	#	Question-Answering	Dataset):	A	Story Cloze Test:	In-	ference	(NLI):							
sen-	=	plex-	2^{H(s)}									read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
tence	-	ity	&=	guage								read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
with	\sum_{i=1}^N	2^{-	model									read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
\$N\$	P(w_i)	the	\frac{1}{s_N}	ld								read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
words,	\log_2	sen-	\sum_{i=1}^N									read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
\$s	p(w_i)	tence	\log_2	dict								read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
=	=	be-	p(w_i)	high								read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
(w_1,	-	comes:	=	word								read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
\dots,	\sum_{i=1}^N	(2^{\sum_{i=1}^N										read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
w_N)\$,	\frac{1}{N}	\log_2	a-									read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
the	\log_2	p(w_i)	\frac{1}{N}									read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
en-	p(w_i)	\frac{1}{N}										read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
tropy	\$\$	=	ties.									read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
looks		(p(w_1)	There-									read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
as		\dots	fore,									read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
fol-		p(w_N))	the									read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
lows,		\frac{1}{s_N}	the									read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
sim-		\end{aligned}										read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
ply	\$\$	plex-										read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
as-		ity										read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
sum-		the										read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
ing		bet-										read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
that		ter.										read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
each												read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
word												read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
has												read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
the												read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
same												read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
fre-												read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
quency,												read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,
\$\frac{1}{N}\$:												read-	ing	com-	mon-	sense	also	known	as	Text	En-	tail-	ment,

Named Entity Recognition (NER)	Sentiment Analysis	Semantic Role Labeling (SRL)	Sentence Classification	Sentence Pair Classification	CoLA
RTE (Recognizing Textual Entailment): A set of words in a text which are the names of things, such as persons and locations, for many names, or general and proper names (reference): A collection of 570k human-written English sentence pairs manually	CoNLL-2003 NLP task: consists of newswire from the Reuters, concentrated on four types of named entities: persons, locations, organizations and names of miscellaneous entities.	SSL (Stanford Sentiment Treebank) the predicate-argument structure dataset of a movie reference, and with bi-ten sentences described as class-si-fi-ing “Who did la-bels to whom”.	CoNLL-2004 & larger CityL-2005 also known as <i>paraphrase detection</i>	MIT RCNP (Multiple Choice Paraphrase task): It contains sentences for grammatical as-pects on ability. web, with annotations indicating whether each pair is semantically equivalent.	CoLA (Corpus of Linguistic Acceptability): a binary single-sentence classification task.
SNLI (Stanford Natural Language Inference): A collection of 570k human-written English sentence pairs manually	OntoNotes 5.0: This corpus contains			QQP (Quora Question Pairs) STS Benchmark:	

e-3e-3	e-3e-3	e-3e-3	e-3e-3	e-3e-3	e-3e-3	e-3e-3	e-3e-3
Text	Part-	Machine	Coreference	Long-	Multi-		
Chunk-	of-	Trans-	Res-	range	task		
ing:	CoNLL-	la-	WES-	CoNLL-	LAMBADA		GLUE
To	2000	tion:	2015	20De-	(Lamb-		multi-
di-	Speech	See	English-	pend-	gung-		task
vide	(POS)	Standard	tion:	dency	Mod-		bench-
a	Tag-	NLP	data-		el-		mark:
text	ging:		(Large		ing		https://gluebenchmark.com/
in	tag	page.	men-		Broad-		
syn-	parts		tions		ened		
tac-	of		WMT		to		
ti-	speech		2014		Ac-		decaNLP
cally	to		text		count		ben-
cor-	each		that		for		mark:
re-	to-		German		Dis-		https://dnnlms.github.io/
lated	ken,		data		course		
parts	such		(Medium)		As-		
of	as		the		pects):		
words.	noun,		IWSLT		A		
	verb,		2015		col-		
	ad-		der-		lec-		
	jec-		English-		tion		
	tive,		Vietnamese		of		
	etc.		data		nar-		
	the		(Small)		ra-		
	Wall		en-		tive		
	Street		ti-		pas-		
	Jour-		ties.		sages		
	nal				ex-		
	por-				tracted		
	tion				from		
	of				the		
	the				Book-		
	Penn				Cor-		
	Tree-				pus		
	bank				and		
	(Mar-				task		
	cus				is		
	et				to		
	al.,				pre-		
	1993).				dict		
					the		
					last		
					word,		
					which		
					re-		
					quire		
					at		
					least		
					50		
					to-		
					kens		
					of		

[illegible]

e-3e-3[6]	e-3e-3[7]	e-3e-3[8]	e-3e-3[9]	e-3e-3[10]	e-3e-3[11]	e-3e-3[12]	e-3e-3[13]	e-3e-3[14]	e-3e-3[15]	e-3e-3[16]
Jeremy Alec Howard Radford Sebastian Ruder. “Universal language model fine-tuning for text classification.”	Alec Radford, Detlev Schreiner, and Karan Aggarwal. “Improving language modeling with dropout.”	Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Christopher D. Manning. “BERT: pre-training of deep bidirectional transformers for language understanding.”	Mike Schuster, Noam Shazeer, and Quoc V. Le. “A neural machine translation system for Google Translate.”	Google’s Neural Machine Translation System: Bridging the Gap between Human and Machine Translation	Ashish Vaswani, Noam Shazeer, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. “Attention is all you need.”	Peter J. Liu, Jiayuan Gu, and Jurgen Schmidhuber. “Generating text by summing long sequences.”	Sebastian Ruder. “Efficient subsequence sampling for neural machine translation.”	Alec Radford, Alec Radford, and Alec Radford. “Language models as unsupervised multitask learners.”	Rico Senrich, Rico Senrich, and Rico Senrich. “Neural machine translation with rare words.”	Zhenzhong Lan, Zhenzhong Lan, and Zhenzhong Lan. “ALBERT: A Lite BERT for Self-supervised Learning of language representations.”
2018. ACL	2018. ACL	2018. ACL	2018. ACL	2018. ACL	2017. NIPS	2018. ICLR	2018. ICLR	2018. ICLR	2015. arXiv	2019. arXiv

Year	Author(s)	Title	Conference/Journal
2017	Yinhan Liu, et al.	“RoBERTa: A Robustly Optimized BERT Pre-training Approach.”	arXiv Preprint arXiv:1907.11692 (2019).
2018	Tom B. Brown, et al.	“Language Models are Few-Shot Learners”	NeurIPS 2020.
2019	Zhilin Yang, et al.	“XLNet: A Generalized Autoregressive Pre-training for Language Understanding.”	NeurIPS 2019.
2019	Mike Lewis, et al.	“BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation, Summarization, and Comprehension.”	ACL 2020.
2020	Kevin Clark, et al.	“TRA: Exploring Text-to-Text Encoders as Discriminators for Natural Language Generation.”	ICLR 2020.
2021	Colin Raffel, et al.	“Exploring Limits of Encoder-Only Transformers”	JMLR 2020.
2021	Architecture	Deep Neural Networks for Long-Duration Reading Transformer Overfitting?	Object Detection Model Part 4: Fast Detection Models

20 Citation

Please cite this work as:

Weng, Lilian. “Title”. Lil ’Log (May 2025). <https://lilianweng.github.io/posts/2019-01-31-lm/>

Or use the BibTeX citation:

```
@article{weng2025,  
title = {},  
author = {Weng, Lilian},  
journal = {lilianweng.github.io},  
year = {2025}, month = {May},  
url = {}  
}
```

21 References

Tags: [tag1](#) [tag2](#)

[Title](#) [Title](#)

Share this article: [Twitter](#) | [LinkedIn](#) | [Reddit](#) | [Facebook](#) | [WhatsApp](#) | [Telegram](#)